

Machine Learning with R at LRZ: Introduction to mlr

Spam E-mail Database

Description

A data set collected at Hewlett-Packard Labs, that classifies 4601 e-mails as spam or non-spam. In addition to this class label there are 57 variables indicating the frequency of certain words and characters in the e-mail.

Format

A data frame with 4601 observations and 58 variables.

The first 48 variables contain the frequency of the variable name (e.g., business) in the e-mail. If the variable name starts with num (e.g., num650) then it indicates the frequency of the corresponding number (e.g., 650). The variables 49-54 indicate the frequency of the characters ‘;’, ‘(’, ‘[’, ‘!’, ‘\\$', and ‘\#’. The variables 55-57 contain the average, longest and total run-length of capital letters. Variable 58 indicates the type of the mail and is either **"nonspam"** or **"spam"**, i.e. unsolicited commercial e-mail.

Details

The data set contains 2788 e-mails classified as **"nonspam"** and 1813 classified as **"spam"**.

The “spam” concept is diverse: advertisements for products/web sites, make money fast schemes, chain letters, pornography... This collection of spam e-mails came from the collectors’ postmaster and individuals who had filed spam. The collection of non-spam e-mails came from filed work and personal e-mails, and hence the word ‘george’ and the area code ‘650’ are indicators of non-spam. These are useful when constructing a personalized spam filter. One would either have to blind such non-spam indicators or get a very wide collection of non-spam to generate a general purpose spam filter.

Source

- Creators: Mark Hopkins, Erik Reeber, George Forman, Jaap Suermondt at Hewlett-Packard Labs, 1501 Page Mill Rd., Palo Alto, CA 94304
- Donor: George Forman (gforman at nospam hpl.hp.com) 650-857-7835

These data have been taken from the UCI Repository Of Machine Learning Databases at <http://www.ics.uci.edu/~mllearn/MLRepository.html>

References

T. Hastie, R. Tibshirani, J.H. Friedman. *The Elements of Statistical Learning*. Springer, 2001.

Exercise

a) Create a binary classification task from the spam data

```
library(mlr)
```

```
## Loading required package: ParamHelpers
```

```
data(spam, package = "kernlab")
```

```
spam.task = makeClassifTask(id = "spam", data = spam, target = "type", positive = "spam")
spam.task
```

```
## Supervised task: spam
```

```
## Type: classif
```

```
## Target: type
```

```
## Observations: 4601
```

```
## Features:
```

```
##      numerics      factors      ordered functionals
```

```
##           57           0           0           0
```

```
## Missings: FALSE
```

```
## Has weights: FALSE
```

```
## Has blocking: FALSE
```

```
## Has coordinates: FALSE
```

```
## Classes: 2
```

```
## nonspam      spam
```

```
##      2788      1813
```

```
## Positive class: spam
```

b) List all learners that could be trained on spam.task

```
listLearners(spam.task, warn.missing.packages = FALSE)
```

```
##              class              name short.name
## 1      classif.ada      ada Boosting      ada
## 2  classif.adaboostm1      ada Boosting M1  adaboostm1
## 3  classif.bartMachine Bayesian Additive Regression Trees bartmachine
## 4      classif.binomial      Binomial Regression      binomial
## 5      classif.boosting      Adabag Boosting      adabag
## 6      classif.bst      Gradient Boosting      bst
##      package      type installed numerics factors ordered missings weights
## 1      ada,rpart classif      TRUE      TRUE      TRUE      FALSE      FALSE      FALSE
## 2      RWeka      classif      TRUE      TRUE      TRUE      FALSE      FALSE      FALSE
## 3  bartMachine      classif      TRUE      TRUE      TRUE      FALSE      TRUE      FALSE
## 4      stats      classif      TRUE      TRUE      TRUE      FALSE      FALSE      TRUE
## 5  adabag,rpart      classif      TRUE      TRUE      TRUE      FALSE      TRUE      FALSE
## 6      bst,rpart      classif      TRUE      TRUE      FALSE      FALSE      FALSE      FALSE
##      prob      oneclass      twoclass      multiclass      class.weights      featimp      oobpreds
## 1      TRUE      FALSE      TRUE      FALSE      FALSE      FALSE      FALSE
## 2      TRUE      FALSE      TRUE      TRUE      FALSE      FALSE      FALSE
## 3      TRUE      FALSE      TRUE      FALSE      FALSE      FALSE      FALSE
## 4      TRUE      FALSE      TRUE      FALSE      FALSE      FALSE      FALSE
## 5      TRUE      FALSE      TRUE      TRUE      FALSE      TRUE      FALSE
## 6      FALSE      FALSE      TRUE      FALSE      FALSE      FALSE      FALSE
##      functionals      single.functional      se      lcens      rcens      icens
## 1      FALSE      FALSE      FALSE      FALSE      FALSE      FALSE
## 2      FALSE      FALSE      FALSE      FALSE      FALSE      FALSE
## 3      FALSE      FALSE      FALSE      FALSE      FALSE      FALSE
```

```
## 4      FALSE      FALSE FALSE FALSE FALSE FALSE
## 5      FALSE      FALSE FALSE FALSE FALSE FALSE
## 6      FALSE      FALSE FALSE FALSE FALSE FALSE
## ... (#rows: 78, #cols: 24)
```

c) Select a learner you like and create it. If you want to can change its hyperparameters

```
lrn = makeLearner("classif.rpart", predict.type = "prob")
```

d) Create an index set of train and test indicies. The test set should have 1000 observations.

d*) Ensure that the fraction between "spam" and "nonspam" is the training and test set is the same as in the full dataset.

```
n = getTaskSize(spam.task)
test.inds = sample(1:n, size = 1000)
train.inds = setdiff(1:n, test.inds)
head(test.inds)
```

```
## [1] 4467 2970 740 1431 3859 515
```

```
head(train.inds)
```

```
## [1] 1 2 4 5 6 7
```

e) Train your model on the train subset of the spam data and predict on the test subset.

```
mod = train(lrn, spam.task, subset = train.inds)
preds = predict(mod, spam.task, subset = test.inds)
print(mod)
```

```
## Model for learner.id=classif.rpart; learner.class=classif.rpart
## Trained on: task.id = spam; obs = 3601; features = 57
## Hyperparameters: xval=0
```

```
print(preds)
```

```
## Prediction: 1000 observations
## predict.type: prob
## threshold: nonspam=0.50,spam=0.50
## time: 0.01
##      id  truth prob.nonspam prob.spam response
## 4467 4467 nonspam    0.92990    0.0701  nonspam
## 2970 2970 nonspam    0.92990    0.0701  nonspam
## 740  740  spam     0.16552    0.8345   spam
## 1431 1431  spam     0.19799    0.8020   spam
## 3859 3859 nonspam    0.92990    0.0701  nonspam
## 515  515  spam     0.04021    0.9598   spam
## ... (#rows: 1000, #cols: 5)
```

f) Evaluate the performance of your model based on accuracy and area under the curve.

```
perf = performance(preds, measures = list(acc, auc))
perf
```

```
##      acc      auc
## 0.8850 0.8992
```

g) Try to find a model with an AUC of at least 98%.

- Try different models
- Change hyperparameters

- Have a closer look at the feature and try to find transformations or combination of features that improve your model's performance

```
lrn2 = makeLearner("classif.randomForest", predict.type = "prob")
mod = train(lrn2, spam.task, subset = train.inds)
preds = predict(mod, spam.task, subset = test.inds)
performance(preds, measures = auc)
```

```
##      auc
## 0.9846
```